

ZELIN TAN

tanzl@mail.ustc.edu.cn · [Homepage](#) · [Google Scholar](#)

EDUCATION

University of Science and Technology of China, Artificial Intelligence, <i>Ph.D. Student</i>	2025.6 - Now
University of Science and Technology of China, Computer Technology, <i>M.Eng.</i>	2022.9 - 2025.6
University of Science and Technology of China, Computer Science <i>B.Eng.</i>	2018.9 - 2022.6

INTERNSHIP

Shanghai AI Lab: Agent&LLM Post Training Mentors: Chen Zhang and Shuyue Hu	2025.7 - Now
• Skill-use training and agent routing: Built verified long-horizon task/trajectory synthesis for SKT and developed the Select-then-Solve reasoning-paradigm router. • RL post-training: Developed PAPO's decoupled advantage normalization and conducted an empirical scaling-law study across model sizes and training regimes. • Agent systems: Contributed to Agents-A1 for long-horizon interaction and AgentPanel for human-AI scientific collaboration.	
Meituan: LLM Inference Optimization Mentors: Yineng Zhang and Ke Bao	2024.5 - 2024.7
• Implemented block-granularity prefix caching and TensorRT operator fusion for production LLM and generative-recommendation inference.	
Tencent: LLM Inference Optimization Mentor: Hande Dong	2023.12 - 2024.5
• Built chunked prefill, GPU-memory-aware dynamic chunk sizing, and SM-aware prefill/decode operator parallelism in an internal vLLM-based serving engine.	

PUBLICATIONS

-
- [SKT: Skill-Use Training at Scale via Verified Synthetic Data Generation](#) (**First author, AAAI 2027 under review**)
 - [PAPO: Stabilizing Rubric Integration Training via Decoupled Advantage Normalization](#) (**First author, EMNLP 2026 under review**)
 - [Scaling Behaviors of LLM Reinforcement Learning Post-Training: An Empirical Study in Mathematical Reasoning](#) (**First author, ACL 2026 Main**)
 - [AgentPanel: Toward a New Paradigm for Human-AI Collaboration in Exploring Scientific Questions](#)
 - [Scaling the Horizon, Not the Parameters: Reaching Trillion-Parameter Performance with a 35B Agent](#)
 - [Parametric Skills](#)
 - [Select-then-Solve: Paradigm Routing as Inference-Time Optimization for LLM Agents](#) (**EMNLP 2026 under review**)
 - [SkillsBench: Benchmarking How Well Agent Skills Work Across Diverse Tasks](#) (**NeurIPS 2026 under review**)
 - [AgentAsk: Multi-Agent Systems Need to Ask](#) (**ACL 2026 Main**)
 - [The Landscape of Agentic Reinforcement Learning for LLMs: A Survey](#) (**TMLR 2025**)

OPEN SOURCE

[ROLL](#) (contributor) · [AgentPanel](#) (contributor) · [SkillsBench](#) (contributor) · [libsmctrl](#) (designer)